

# BIOSTATS

Akira Terui

Last updated on 2022-10-31



# Contents

|   |           |
|---|-----------|
| <b>Introduction</b>                               | <b>5</b>  |
| <b>1 Project Management</b>                       | <b>7</b>  |
| 1.1 R Project . . . . .                           | 7         |
| 1.2 Robust coding . . . . .                       | 10        |
| <b>2 Why statistics?</b>                          | <b>11</b> |
| 2.1 Rationale . . . . .                           | 11        |
| 2.2 Example – fish density in a lake . . . . .    | 11        |
| 2.3 Mean & Variance . . . . .                     | 13        |
| 2.4 Probability distribution . . . . .            | 14        |
| <b>3 Appendix: <i>Git</i> &amp; <i>GitHub</i></b> | <b>15</b> |
| 3.1 <i>Git</i> & <i>GitHub</i> . . . . .          | 15        |
| 3.2 Commit & Push . . . . .                       | 17        |
| <b>4 Appendix: Data Structure</b>                 | <b>25</b> |
| 4.1 Overview . . . . .                            | 25        |
| 4.2 Vector . . . . .                              | 25        |
| 4.3 Matrix . . . . .                              | 29        |
| 4.4 Data Frame . . . . .                          | 32        |
| 4.5 Exercise . . . . .                            | 33        |



# Introduction

This textbook aims to introduce fundamental statistical techniques and their applications to biological data. A unique aspect of this book is the “flipped-order” introduction. Many statistics courses start with theory; yet, I found it difficult for those unfamiliar with statistics. I will start with a real example of the method, followed by the explanation for an underlying theory/concept. The author is an ecologist, so some methods in this book might not be popular in other fields.



# Chapter 1

## Project Management

### 1.1 R Project

For this entire book, I will use **R Project** as a fundamental unit of work-space, in which all the relevant materials (e.g., R scripts `.R` and data files) are assembled together. There are many ways to organize your project, but I usually make a single **R Project** for a collection of scripts and data that will lead to a single publication (see example here). To setup an **R Project** you will need *RStudio* in addition to base *R*. While *R* is a stand-alone software, I strongly recommend to use it with *RStudio*. *RStudio* has many functions that help your data analysis. *R* and *RStudio* can be installed from the following websites:

- R (you can choose any CRAN mirror to download)
- RStudio

Once you open *RStudio*, you will see the following interface (Figure 1.1). There are three major panels in its first appearance – **Console**, **Environment**, and **Files**. **Console** is the place where you write your codes and execute calculation/data manipulation/analysis. **Environment** lists items you saved as an object. **Files** list any files in a designated location in your computer.

Let's type something in console to see what happens. Copy and paste the following script to **Console**. You should see `x` in the environment.

```
x <- c(1, 2)
```

`x` is an **object** where information is saved - in this case, we stored a sequence of 1 and 2 in object `x`. Once you store the information into `x`, you can call it just typing `x`.

```
x
```

```
## [1] 1 2
```

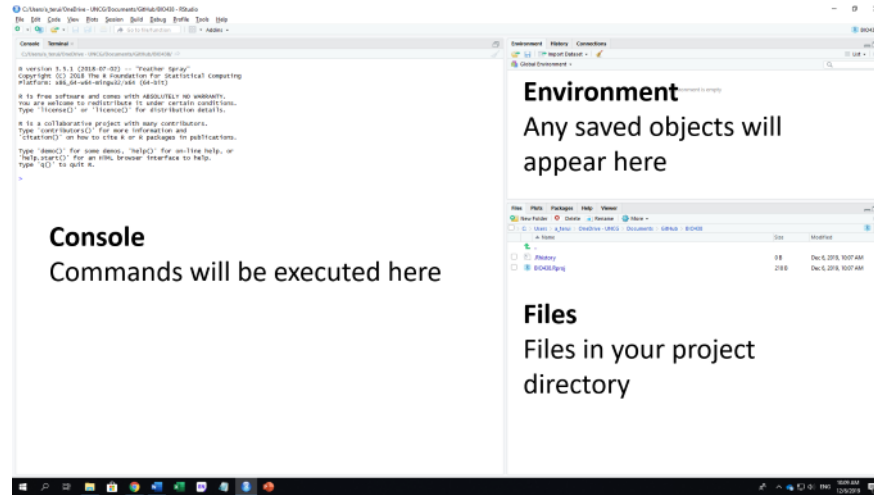


Figure 1.1: RStudio interface.

Sweet! You can work on your data like this, however, **it is actually a BAD idea**. As you work on your project, numerous materials will be generated (e.g., I would write > 2000 lines of codes for one project). How do you manage your codes?

### 1.1.1 Script Editor

It is highly recommended to manage your script in **Editor** instead. **Editor** is where you draft and tweak your codes before executing them in **Console**. To create space for **Editor**, hit **Ctrl + Shift + N**. A new panel pops up on the top left. Let's type the following script in **Editor**.

```
y <- c(3, 4)
```

Then, hit **Ctrl + S** to save the **Editor** file. *RStudio* will prompt you to enter the file name of the **Editor**<sup>1</sup>.

### 1.1.2 File Name

It is also critical to have **consistent naming rules** for your files. As you make progress on your project, the number of files in each sub-directory will increase, perhaps exponentially. You will find it difficult navigating yourself unless you have clear naming rules for files (and even worse for others). Here are some recommendations:

- **NO SPACE**. Use underscore.
  - Do: `script_week1.R`

<sup>1</sup>In *R*, an editor file has an extension of `.R`.



- Don't: `script week1.R`
- **NO UPPERCASE.** Use lowercase for file names.
  - Do: `script_week1.R`
  - Don't: `Script_week1.R`
- **BE CONSISTENT.** Apply consistent naming rules within a project.
  - Do: R scripts for figures always start with a common prefix, e.g., `figure_XXX.R` `figure_YYY.R` (XXX and YYY specifies further details).
  - Don't: R scripts for figures start with random text, e.g., `XXX_fig.R`, `Figure_Y2.R`, `plotB.R`.

### 1.1.3 Structure Your Project

If you do not save or randomly locate your codes within your computer, you will lose necessary items sooner or later. For this reason, I assemble all the relevant materials in a single R Project. You can create a new R Project with the following procedure.

- a. Go to **File > New Project** on the top menu
- b. Select **New Directory**
- c. Select **New Project**

A new window pops up and prompts you to name a directory with a location in your computer. Click **Browse** to select a location for the directory<sup>2</sup>.

The internal structure of an R Project is extremely important to navigate yourself (others once it's published). R Project will be composed of multiple types of files, typically `.R`, `.csv`, `.rds`, `.Rmd` among others. Unless those files are arranged in an organized manner, it is VERY LIKELY to make severe errors in coding. So I take this seriously. Table 1.1 is my suggested subdirectory structure.

Table 1.1: Suggested internal structure of R Project

| Name                      | Content   |
|---------------------------|---|
| <code>README.md</code>    | Markdown file explaining contents in the R Project. Can be derived from <code>README.Rmd</code> .   |
| <code>/code</code>        | Sub-directory for R scripts ( <code>.R</code> ).  |
| <code>/data_raw</code>    | Sub-directory for raw data before data manipulation ( <code>.csv</code> or other formats). Files in this sub-directory MUST NOT be modified unless there are changes to raw data entries. |
| <code>/data_format</code> | Sub-directory for formatted data ( <code>.csv</code> , <code>.rds</code> , or other formats).   |
| <code>/output</code>      | Sub-directory for result outputs ( <code>.csv</code> , <code>.rds</code> , or other formats). This may include statistical estimates from linear regression models etc.                   |

<sup>2</sup>When you locate your project directories in your computer, I would strongly recommend to create a designated space. For example, in my computer, I have a folder named `/r_project` in which all R Project directories are located.

---

| Name              | Content  |
|-------------------|--|
| <code>/rmd</code> | (Optional) Sub-directory for Rmarkdown files ( <code>.Rmd</code> ). Rmarkdown allows seamless integration of R scripts and text. |

---

## 1.2 Robust coding

While this is not mandatory, I strongly recommend to use *RStudio* with *Git* & *GitHub* (see Chapter ?? Appendix for guidance). Coding is innately error-prone<sup>3</sup>; every programmers, even bright ones, make mistakes - no exception. However, the critical difference between beginner and advanced programmers is whether they develop robust coding rules with a self-error-detection system. *Git* is the key to this. I will touch on this occasionally throughout this book.

---

<sup>3</sup>well WAY BETTER than manual data manipulation!

## Chapter 2

# Why statistics?

### 2.1 Rationale

*Why do I need statistics in ecology in the first place?* - this is the first question I got when I stepped into the field of ecology. I thought this would be an easy question that someone could give me the answer right away, but actually it is a deep one. The short answer is “*we need statistics because we only have **partial** information about what we want to know.*” The long answer? See below.

### 2.2 Example – fish density in a lake

Let say we are interested in fish density in a lake (total area  $15 \times 15 = 225 \text{ m}^2$ ). We try to quantify it by catching fish in a certain area ( $1 \text{ m}^2$  for example). Since there should be some variation in fish density among sampled points, we decided to repeat this 10 times to take an average. Figure 2.1 shows a graphical representation of this example.

Through this sampling, we got the following dataset:

```
## # A tibble: 10 x 2
##   plot count
##   <int> <int>
## 1     1     1
## 2     2     1
## 3     3     1
## 4     4     2
## 5     5     7
## 6     6     2
## 7     7     5
## 8     8     1
```




Figure 2.1: Lake fish density example. Points are fish, and white boxes are sampling sites with 1 sq-m

```
## 9      9      3
## 10     10     3
```

Based on this dataset, we may conclude the average fish density in this lake is 2.6. However, how confident are you on this conclusion? The lake area is  $15 \times 15 = 225 \text{ m}^2$  but we sampled only  $1 \text{ m}^2 \times 10 = 10 \text{ m}^2$ ! Indeed, we missed a lot of fish in the lake (Figure 2.2).




Figure 2.2: True fish distribution in the lake

In ecological research (and almost all research field), we are rarely able to obtain the perfect dataset with no uncertainty – the data are partial information of what we want to know (true fish density in the entire lake = 2.22). This is the reason why we refer to a data point as a **sample**, and the average fish density in white boxes (2.6) is the **estimate** of the true fish density because it is the inference from the incomplete information. Since the average fish density is the estimate, it does not perfectly match the true density – how do we account for this uncertainty?

## 2.3 Mean & Variance

In the above section, I talked only about the **mean** (total fish counts / number of plots = mean density per plot area [ $\text{m}^2$ ]). Let me re-write this in a little more mathematical form – I write a single data point (fish count in a single white box) as  $y_i$  with  $i$  representing an ID of a plot (therefore,  $y_1 = 1$ ,  $y_2 = 1$ ,  $y_3 =$

1,  $y_4 = 2, \dots, y_{10} = 3$ ), and refer to the number of plots as  $N$ . The mean fish counts is:

$$\frac{\sum_{i=1}^N y_i}{N}$$

However, the data contain another important aspect of information: variability in fish counts among samples. This variability – called **variance** – expresses the degree of uncertainty in your data. This uncertainty is can be calculated from the deviation of each data point  $x_i$  from the mean  $\mu$ :

## 2.4 Probability distribution

## Chapter 3

# Appendix: *Git* & *GitHub*

### 3.1 *Git* & *GitHub*

In this section, I will cover how to integrate *Git* and *GitHub* into *R Studio*. *R Studio* is excellent as is, but it becomes even better when combined with *Git* and *GitHub*. *Git* is a free and open source distributed version control system. ***Git* tracks changes in your codes codes while you work on your project so you are aware of any changes you made to your script (and other) files.** Tracking changes is extremely important to avoid unintended errors in your codes. This feature also helps avoid creating redundant files. While *Git* is a local system, it has an online counterpart called *GitHub*.

To make this system work, you'll need to go through some processes. The first step is to install *Git* onto your computer:

- **Windows:** Install *Git* from [here](#). You will be asked about “Adjusting your PATH environment”. Choose “*Git* from the command line and also from 3rd-party software” if it is not selected.
- **Mac:** Follow the instruction [here](#).

Then open R Studio and do **Create Project > New Directory > New Project**. If you see a check box **Create a git repository**, check and create a new project (Figure 3.1). You will see a *Git* pane on the upper right panel.

If you can't find the above, do the following: **Tools** in the menu bar > **Terminal > New Terminal**, and type **where git** in the terminal. This will tell you where git executable is located in your computer. Then, go to **Tools** in the menu bar > **Global Options > Git/SVN**. You will see *Git executable* in the box, where you can specify the location of git executable.

Next, go to *GitHub* and create your account (free!). But, give some thoughts on your user name. My advice is the following. First, **use lowercase only**.




Figure 3.1: After installing Git, you should see ‘Create a git repository’.



Second, **include your name** to make it easy to find you on *GitHub*.

*R Studio* works seamlessly with *Git* or *GitHub*, but it is helpful to use a *Git client* as it provides visual aids. There are choices for a *Git client* (see options here) but I will use *GitHub Desktop* (available here) for our exercise. Install GitHub Desktop onto your computer.

## 3.2 Commit & Push

### 3.2.1 Register Your Git repo

Open the *R Project* you've just created as a git repository. Let's make a sample *.R* file (Ctrl + Shift + N) and save it (name as **sample.R**). For example:

```
## produce 100 random numbers that follows a normal distribution
x <- rnorm(100, mean = 0, sd = 1)

## estimate mean
mean(x)

## estimate SD
sd(x)
```

Open GitHub Desktop App. You will see the following GUI (Figure 3.2):

Hit **current repository** (top left) and **Add > Add existing repository** (Figure 3.3):

### 3.2.2 Commit

GitHub Desktop will prompt you to enter a local path to your *Git* repository. Browse and select your directory of the *R Project* - the local *Git* repository will show up in the list of repositories in *GitHub Desktop* (left side bar in Figure 3.3). Now, you are ready to **Commit** your files to *Git*! Commit is the procedure to record your file change history in *Git*. To make this happen, click your *Git* repository on GitHub Desktop, and you will see the following:

At this stage, your file (*.R* or else) is saved onto your local directory, **but has not been recorded in *Git* as a change history**. To make a Commit, the first thing to do is to choose a file(s). There are checkboxes next to each of the new files. If this box is checked, you are going to commit changes to *Git*. Once you selected the files you want to make a commit, go to the bottom left of the window. There is a small box saying **summary (required)** or **Create sample.R**. This is the place where you can put any title that describes what you did in this commit, and **you cannot Commit unless you enter this information!** For example, I would write **initial commit** for this exercise - from the second commit, you should put a more informative commit message so you can track changes when needed. You can google recommendations for

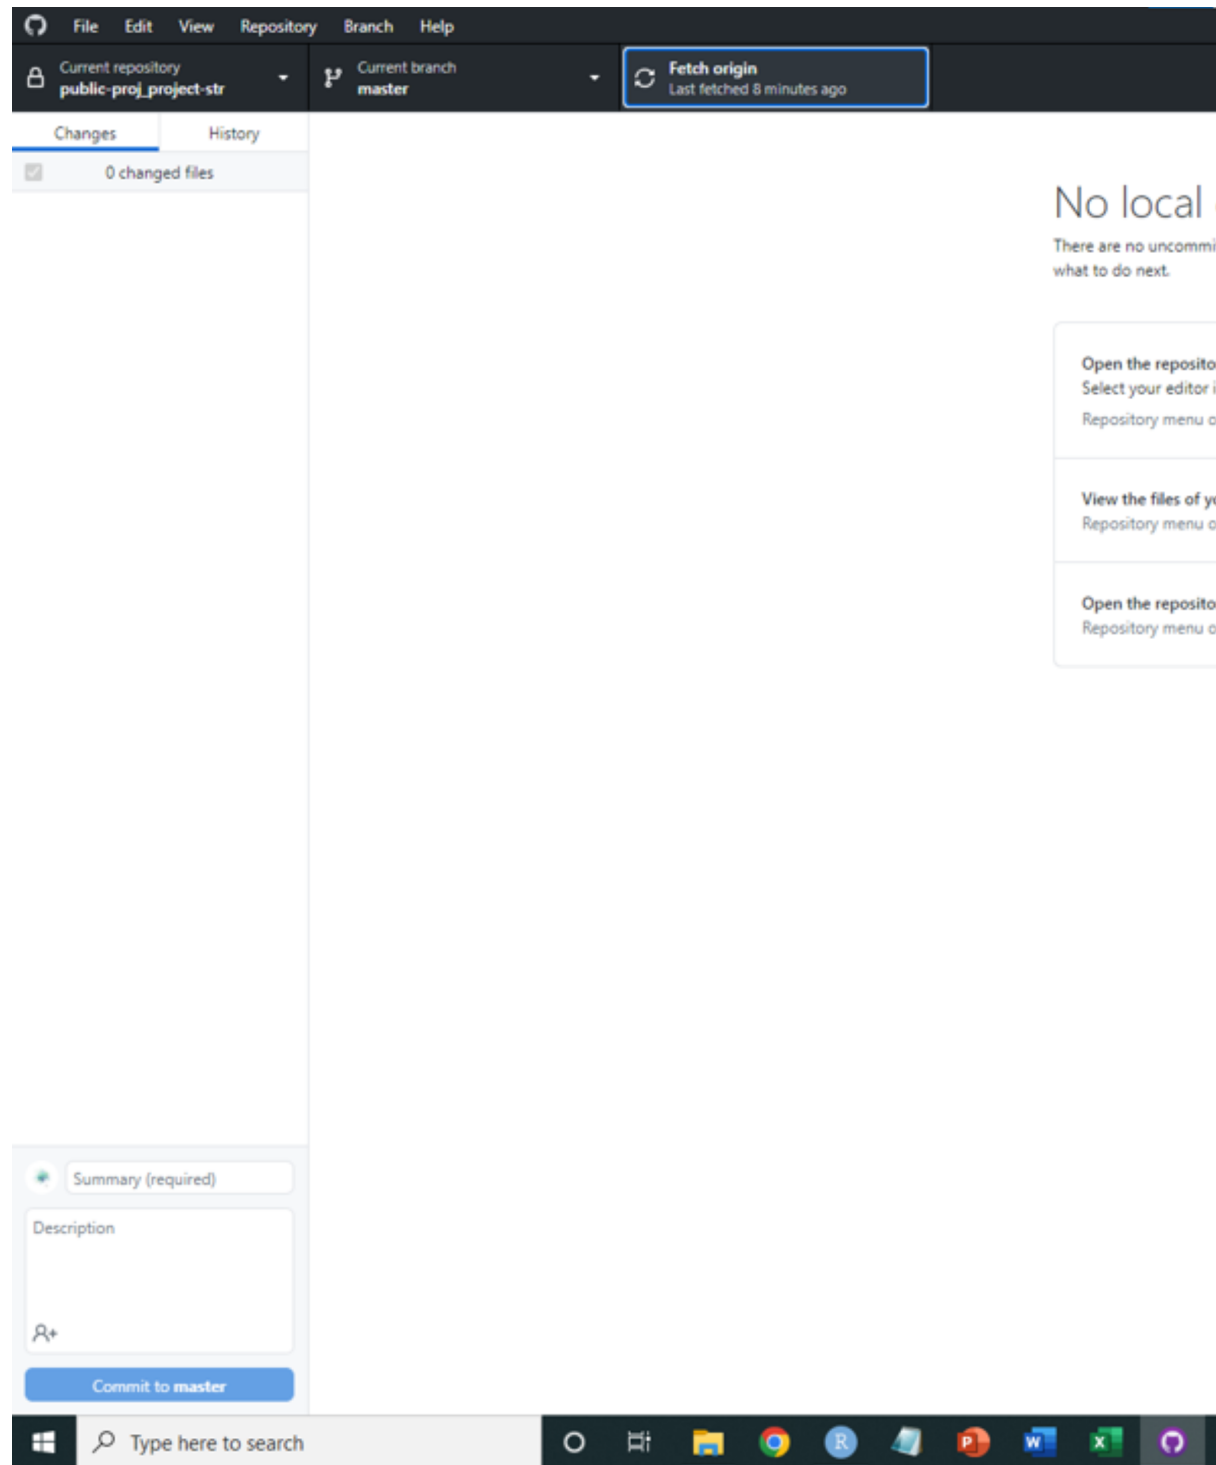


Figure 3.2: GUI for GitHub Desktop




Figure 3.3: Add dropdown on the top left




Figure 3.4: You will see ‘Create sample.R’ or ‘summary (required)’ on the bottom left

how to write commit titles/descriptions. Then hit `commit to master`. **Now, changes to the selected files have been recorded in *Git*!**

### 3.2.3 Push

Remember, your changes are recorded in your local computer **but not published in your online repository!** To send local changes to the online *GitHub* repository, you will need to **Push** commits via *GitHub Desktop*. Push is the procedure to send your Commit(s) to *GitHub*. Once you do a Commit, *GitHub Desktop* will ask you whether you want to Push it to an online repository (Figure 3.5). If this is the first push, there is no corresponding repository on *GitHub* tied to your local repository, so *GitHub Desktop* will ask you if you want to publish it on *GitHub* (NOTE: although it says ‘publish’, your repository will remain private unless you explicitly tell *Git Desktop* to make it public). If you are comfortable with the changes you made, **Push** it!

### 3.2.4 Edit

We went through how we get things uploaded onto *GitHub*, but what happens if we make changes to existing files? To see this, let’s make a minor change to your R script. We have created a file `sample.R`:

```
## produce 100 random numbers that follows a normal distribution
x <- rnorm(100, mean = 0, sd = 1)

## estimate mean
mean(x)

## estimate SD
sd(x)
```

Edit this as follows:

```
## produce 100 random numbers that follows a normal distribution
x <- rnorm(100, mean = 0, sd = 1)

## estimate mean
median(x)

## estimate SD
var(x)
```

After making this change, go to *GitHub Desktop* again. *GitHub Desktop* automatically detects the difference between the new and old files and shows which part of the script has been edited! This helps coding quite a bit:




Figure 3.5: To Push your code, hit the highlighted menu button




Figure 3.6: Git detects edits to your codes





## Chapter 4

# Appendix: Data Structure

### 4.1 Overview

R has 6 basic **data types**.

- character: "aquatic", "ecology" (no order)
- factor: similar to character, but has *levels* (alphabetically ordered by default)
- numeric: 20.0 , 15.5
- integer: 3, 7
- logical: TRUE , FALSE
- complex: 1+2i (complex numbers with real and imaginary parts)

These elements form one of the following **data structures**.

- **vector**: a series of elements. A single data type is allowed in a single vector
- **matrix**: elements organized into rows and columns. A single data type is allowed in a single matrix
- **data frame**: looks similar to a matrix, but allows different data types in different columns

### 4.2 Vector

#### 4.2.1 Create Vector

Below are examples of atomic character vectors, numeric vectors, integer vectors, etc. There are many ways to create vector data. The following examples use `c()`, `:`, `seq()`, `rep()`:

```
#ex.1a manually create a vector using c()
```

```
x <- c(1,3,4,8)
```

```
x
```

```
## [1] 1 3 4 8
```

```
#ex.1b character
```

```
x <- c("a", "b", "c")
```

```
x
```

```
## [1] "a" "b" "c"
```

```
#ex.1c logical
```

```
x <- c(TRUE, FALSE, FALSE)
```

```
x
```

```
## [1] TRUE FALSE FALSE
```

```
#ex.2 sequence of numbers
```

```
x <- 1:5
```

```
x
```

```
## [1] 1 2 3 4 5
```

```
#ex.3a replicate same numbers or characters
```

```
x <- rep(2, 5) # replicate 2 five times
```

```
x
```

```
## [1] 2 2 2 2 2
```

```
#ex.3b replicate same numbers or characters
```

```
x <- rep("a", 5) # replicate "a" five times
```

```
x
```

```
## [1] "a" "a" "a" "a" "a"
```

```
#ex.4a use seq() function
```

```
x <- seq(1, 5, by = 1)
```

```
x
```

```
## [1] 1 2 3 4 5
```

```
#ex.4b use seq() function
```

```
x <- seq(1, 5, by = 0.1)
```

```
x
```

```
## [1] 1.0 1.1 1.2 1.3 1.4 1.5 1.6 1.7 1.8 1.9 2.0 2.1 2.2 2.3 2.4 2.5 2.6 2.7 2.8
```

```
## [20] 2.9 3.0 3.1 3.2 3.3 3.4 3.5 3.6 3.7 3.8 3.9 4.0 4.1 4.2 4.3 4.4 4.5 4.6 4.7
```

```
## [39] 4.8 4.9 5.0
```

### 4.2.2 Check Features

R provides many functions to examine features of vectors and other objects, for example:

- `class()` - return high-level data structure of the object
- `typeof()` - return low-level data structure of the object
- `attributes()` - metadata of the object
- `length()` - number of elements in the object
- `sum()` - sum of object's elements
- `mean()` - mean of object's elements

#### Numeric Vector

```
x <- c(1.2, 3.1, 4.0, 8.2)
x
```

```
## [1] 1.2 3.1 4.0 8.2
```

```
class(x)
```

```
## [1] "numeric"
```

```
typeof(x)
```

```
## [1] "double"
```

```
length(x)
```

```
## [1] 4
```

```
sum(x)
```

```
## [1] 16.5
```

```
mean(x)
```

```
## [1] 4.125
```

#### Character Vector

```
## [1] "character"
```

```
## [1] 3
```

### 4.2.3 Access

#### Element ID

Use brackets `[]` when accessing specific elements in an object. For example, if you want to access element #2 in the vector `x`, you may specify as `x[2]`:

```
## [1] 2
```

```
## [1] 2 2
```

```
## [1] 2 3 2
```

### Equation

R provides many ways to access elements that suffice specific conditions. You can use mathematical symbols to specify what you need, for example:

- `==` equal
- `>` larger than
- `>=` equal & larger than
- `<` smaller than
- `<=` equal & smaller than
- `which()` a function that returns element `#` that suffices the specified condition

The following examples return a logical vector indicating whether each element in `x` suffices the specified condition:

```
# creating a vector
x <- c(2,2,3,2,5)

# ex.1a equal
x == 2
```

```
## [1] TRUE TRUE FALSE TRUE FALSE

# ex.1b larger than
x > 2
```

```
## [1] FALSE FALSE TRUE FALSE TRUE
```

You can access elements that suffice the specified condition using brackets, for example:

```
# ex.2a equal
x[x == 2]
```

```
## [1] 2 2 2

# ex.2b larger than
x[x > 2]
```

```
## [1] 3 5
```

Using `which()`, you can see which elements (i.e., `#`) matches what you need:

```
# ex.3a equal
which(x == 2) # returns which elements are equal to 2
```

```
## [1] 1 2 4

# ex.3b larger than
which(x > 2)
```

```
## [1] 3 5
```

## 4.3 Matrix

### 4.3.1 Create Matrix

Matrix is a set of elements (*single data type*) that are organized into rows and columns:

```
#ex.1 cbind: combine objects by column
```

```
x <- cbind(c(1,2,3), c(4,5,6))
x
```

```
##      [,1] [,2]
## [1,]    1    4
## [2,]    2    5
## [3,]    3    6
```

```
#ex.2 rbind: combine objects by row
```

```
x <- rbind(c(1,2,3), c(4,5,6))
x
```

```
##      [,1] [,2] [,3]
## [1,]    1    2    3
## [2,]    4    5    6
```

```
#ex.3 matrix: specify elements and the number of rows (nrow) and columns (ncol)
```

```
x <- matrix(1:9, nrow = 3, ncol = 3)
x
```

```
##      [,1] [,2] [,3]
## [1,]    1    4    7
## [2,]    2    5    8
## [3,]    3    6    9
```

### 4.3.2 Check Features

R provides many functions to examine features of matrix data, for example:

- `dim()` number of rows and columns
- `rowSums()` row sums
- `colSums()` column sums

#### Integer Matrix

```
x <- matrix(1:9, nrow = 3, ncol = 3)
x
```

```
##      [,1] [,2] [,3]
## [1,]    1    4    7
## [2,]    2    5    8
## [3,]    3    6    9
```

```

class(x)

## [1] "matrix" "array"
typeof(x)

## [1] "integer"
dim(x)

## [1] 3 3

```

**Character Matrix**

```

y <- matrix(c("a", "b", "c", "d", "e", "f"), nrow = 3, ncol = 2)
y

##      [,1] [,2]
## [1,] "a"  "d"
## [2,] "b"  "e"
## [3,] "c"  "f"
class(y)

## [1] "matrix" "array"
typeof(y)

## [1] "character"
dim(y)

## [1] 3 2

```

### 4.3.3 Access

When accessing matrix elements, you need to pick row(s) and/or column(s), for example:

```

x <- matrix(1:9, nrow = 3, ncol = 3)
x

##      [,1] [,2] [,3]
## [1,]  1   4   7
## [2,]  2   5   8
## [3,]  3   6   9
x[2,3] # access an element in row #2 and column #3

## [1] 8
x[2,] # access elements in row #2

```

```
## [1] 2 5 8
x[c(2,3),] # access elements in rows #2 and 3

##      [,1] [,2] [,3]
## [1,]    2    5    8
## [2,]    3    6    9
x[,c(2,3)] # access elements in columns #2 and 3

##      [,1] [,2]
## [1,]    4    7
## [2,]    5    8
## [3,]    6    9
```

You can assess each element with mathematical expressions just like vectors:

```
x == 2 # equal

##      [,1] [,2] [,3]
## [1,] FALSE FALSE FALSE
## [2,]  TRUE FALSE FALSE
## [3,] FALSE FALSE FALSE
x > 2 # larger than

##      [,1] [,2] [,3]
## [1,] FALSE TRUE TRUE
## [2,] FALSE TRUE TRUE
## [3,]  TRUE TRUE TRUE
```

However, care must be taken when accessing elements, as it will be automatically converted to vector data:

```
## [1] 2
## [1] 3 4 5 6 7 8 9
```

`which()` needs an additional argument to return both row and column #:

```
##      row col
## [1,]    2    1

##      row col
## [1,]    3    1
## [2,]    1    2
## [3,]    2    2
## [4,]    3    2
## [5,]    1    3
## [6,]    2    3
## [7,]    3    3
```

## 4.4 Data Frame

Data frame is a set of elements that are organized into rows and columns, but differ from matrix in several ways.

- it allows *multiple data types* in different columns
- each column has its *name*
- you can access columns by name (using \$)

**Data frame is the most common data structure when manipulating ecological data.** A data set loaded from a spread sheet (we will address this later) will be automatically recognized as a data frame. Here is an example:

### Create a data frame

In the following example, variables `x` and `y` are organized into a single data frame `dat`. Variable are renamed when creating a data frame composed of `x` and `y`.

```
# Create data frame
x <- c("Pristine", "Pristine", "Disturbed", "Disturbed", "Pristine") # Lake type
y <- c(1.2, 2.2, 10.9, 50.0, 3.0) # TSS: total suspended solids (mg/L)
dat <- data.frame(LakeType = x, TSS = y) # x is named as "LakeType" while y is named as "TSS"
dat
```

```
##      LakeType  TSS
## 1  Pristine  1.2
## 2  Pristine  2.2
## 3 Disturbed 10.9
## 4 Disturbed 50.0
## 5  Pristine  3.0
```

### Call column names

```
colnames(dat) # call column names
```

```
## [1] "LakeType" "TSS"
```

### Access by columns

```
dat$LakeType # access LakeType
```

```
## [1] "Pristine" "Pristine" "Disturbed" "Disturbed" "Pristine"
```

```
dat$TSS # access TSS
```

```
## [1] 1.2 2.2 10.9 50.0 3.0
```

You can access elements like a matrix as well:

```
dat[,1] # access column #1
```

```
## [1] "Pristine" "Pristine" "Disturbed" "Disturbed" "Pristine"
```



```
dat[1,] # access row #1

##   LakeType TSS
## 1 Pristine 1.2

dat[c(2,4),] # access row #2 and 4

##   LakeType TSS
## 2 Pristine 2.2
## 4 Disturbed 50.0
```

## 4.5 Exercise

### 4.5.1 Vector

- Create three numeric vectors with length 3, 6 and 20, respectively. Each vector must be created using different functions in R.
- Create three character vectors with length 3, 6 and 20, respectively. Each vector must be created using different functions in R.
- Copy the following script to your R script and perform the following analysis:
  - Identify element IDs of `x` that are greater than 2.0
  - Identify element values of `x` that are greater than 2.0

### 4.5.2 Matrix

- Create a numeric matrix with 4 rows and 4 columns. Each column must contain identical elements.
- Create a numeric matrix with 4 rows and 4 columns. Each row must contain identical elements.
- Create a character matrix with 4 rows and 4 columns. Each column must contain identical elements.
- Create a character matrix with 4 rows and 4 columns. Each row must contain identical elements.
- Copy the following script to your R script and perform the following analysis:
  - Identify element IDs of `x` that are greater than 2.0 (**specify row and column IDs**)
  - Identify element values of `x` that are greater than 2.0 and calculate the mean.

### 4.5.3 Data Frame

- Create a data frame of 3 variables with 10 elements (name variables as `x`, `y` and `z`. `x` must be `character` while `y` and `z` must be `numeric`).

- b. Check the data structure (higher-level) of **x**, **y** and **z**
- c. Copy the following script to your R script and perform the following analysis:
  - Calculate the means of **temperature** and **abundance** for states **VA** and **NC** separately.